

Constraint Satisfaction Problems

Bùi Tiến Lên

01/09/2019



KHOA CÔNG NGHỆ THÔNG TIN
TRƯỜNG ĐẠI HỌC KHOA HỌC TỰ NHIÊN



1. **Defining Constraint Satisfaction Problems**
2. **Backtracking Search for CSPs**
3. **Constraint Propagation: Inference in CSPs**
4. **Local Search for CSPs**
5. **The Structure of Problems**



Defining Constraint Satisfaction Problems

Constraint Satisfaction Problems



- In standard search problem, **state** is a “black box”
- A **constraint satisfaction problem** (CSP) use a **factored representation** for each state.
 - State = a set of variables and each of which has a value
 - Solution = each variable has a value that satisfies all constraints on that variable
- Take advantage of the structure of states
- General-purpose rather than problem-specific heuristics
 - Identify combinations of variable-value that violate the constraints → eliminate large portions of the search space all at once
 - Solutions to complex problems

Constraint Satisfaction Problems (cont.)



A constraint satisfaction problem consists of three components, \mathcal{X} , \mathcal{D} , and \mathcal{C} :

- \mathcal{X} is a set of variables, $\{X_1, \dots, X_n\}$
- \mathcal{D} is a set of domains, $\{D_1, \dots, D_n\}$
 - Each domain D_i consists of a set of allowable values, $\{v_1, \dots, v_k\}$ for variable X_i
- \mathcal{C} is a set of constraints that specify allowable combinations of values.
 - Each constraint C_j consists of a pair $\langle \text{scope}, \text{rel} \rangle$, where *scope* is a tuple of variables that participate in the constraint and *rel* is a relation that defines the values that those variables can take on

State Space And Solution



- Each state in a CSP is defined by an **assignment** of values to **some** or **all** of the variables,


$$\{X_i = v_i, X_j = v_j, \dots\}$$

- An assignment that does not violate any constraints is called a **consistent** or legal assignment.
- A **complete assignment** is one in which every variable is assigned
- A **partial assignment** is one that assigns values to only some of the variables.
- A **solution** to a CSP is a consistent, complete assignment.

Example problem: Map coloring



- The principal states and territories of Australia. Coloring this map can be viewed as a constraint satisfaction problem (CSP). The goal is to assign colors to each region so that no neighboring regions have the same color.
- The map-coloring problem represented as a constraint graph.
- **Constraint graph:** nodes are variables, arcs show constraints
- CSP algorithms use the graph structure to speed up search






Example problem: Map coloring (cont.)

- **Variables:** $\mathcal{X} = \{WA, NT, Q, NSW, V, SA, T\}$
- **Domains:** $D_i = \{red, green, blue\}$
- **Constraints:** *adjacent regions must have different colors*

$$\mathcal{C} = \{SA \neq WA, SA \neq NT, SA \neq Q, SA \neq NSW, SA \neq V, \\ WA \neq NT, NT \neq Q, Q \neq NSW, NSW \neq V\}$$

There are many possible solutions to this problem, e.g.,

$$\{WA = red, NT = green, Q = red, NSW = green, \\ V = red, SA = blue, T = green\}$$



Example problem: Job-shop scheduling



- We consider a small part of the car assembly, consisting of 15 tasks:
 - install axles (front and back)
 - affix all four wheels (right and left, front and back)
 - tighten nuts for each wheel
 - affix hubcaps
 - inspect the final assembly.
- Some tasks must occur before another while many other tasks can go on at once.
 - E.g., a wheel must be installed before the hubcap is put on
- A task takes a certain amount of time to complete.

Example problem: Job-shop scheduling (cont.)



We can represent the tasks with 15 variables:

- **Variables:**

$$\mathcal{X} = \{Axle_F, Axle_B, \\ Wheel_{RF}, Wheel_{LF}, Wheel_{RB}, Wheel_{LB}, \\ Nuts_{RF}, Nuts_{LF}, Nuts_{RB}, Nuts_{LB}, \\ Cap_{RF}, Cap_{LF}, Cap_{RB}, Cap_{LB}, \\ Inspect\}$$

- **Domains:** D_i : the value of each variable is the time that the task starts

Example problem: Job-shop scheduling (cont.)



- **Constraints:** Assume task T_1 and T_2 take duration d_1 and d_2 to complete
 - **Precedence constraints:** task T_1 must occur before task T_2 ;

$$T_1 + d_1 \leq T_2$$

- **Disjunctive constraint:** task T_1 and task T_2 must not overlap in time

$$T_1 + d_1 \leq T_2 \text{ or } T_2 + d_2 \leq T_1$$

Example problem: Job-shop scheduling (cont.)



- The axles have to be in place before the wheels are put on, and it takes 10 minutes to install an axle

$$Axle_F + 10 \leq Wheel_{RF} \quad Axle_F + 10 \leq Wheel_{LF}$$

$$Axle_B + 10 \leq Wheel_{RB} \quad Axle_B + 10 \leq Wheel_{LB}$$

- For each wheel, we must affix the wheel (which takes 1 minute), then tighten the nuts (2 minutes), and finally attach the hubcap (1 minute)

...

- Suppose we have four workers to install wheels, but they have to share one tool that helps put the axle in place

...

Example problem: Job-shop scheduling (cont.)




- The inspection comes last and takes 3 minutes

$$\forall X \neq \textit{Inspect}, X + d_X \leq \textit{Inspect}$$

- Finally, suppose there is a requirement to get the whole assembly done in 30 minutes \rightarrow limit the domain of all variables to

$$D_j = \{1, 2, 3, \dots, 27\}$$

Example problem: 4-Queens



- **Variables:** $\mathcal{X} = \{Q_1, Q_2, Q_3, Q_4\}$
- **Domains:** $D_i = \{1, 2, 3, 4\}$
- **Constraints:**
 - $Q_i \neq Q_j$ (cannot be in same row)
 - $|Q_i - Q_j| \neq |i - j|$ (or same diagonal)



Example problem: Cryptarithmic

Defining
Constraint
Satisfaction
Problems

Backtracking
Search for
CSPs


Constraint
Propagation:
Inference in
CSPs

Local Search
for CSPs

The Structure
of Problems

$$\begin{array}{r} T \ W \ O \\ + \ T \ W \ O \\ \hline F \ O \ U \ R \end{array}$$

(a)



(b)

- **Variables:** $\mathcal{X} = \{F, T, U, W, R, O, C_1, C_2, C_3\}$
- **Domains:** $D_i = \{0, 1, 2, 3, 4, 5, 6, 7, 8, 9\}$
- **Constraints:**
 - $Alldiff(F, T, U, W, R, O)$
 - $T \neq 0, F \neq 0$
 - $C_3 = F, \dots$



Why formulate a problem as a CSP

- Provide natural representation for a wide variety of problems
- Many problems **intractable** in regular state-space search can be solved quickly with CSP formulation.
For example, once we have chosen $\{SA = blue\}$ in the Australia problem.
 - Search: $3^5 = 243$ assignments
 - CSP: $2^5 = 32$ assignments
- Better insights to the problem and its solution

Varieties of CSPs



- Discrete variables and finite domains
 - n variables with size $d \rightarrow O(n^d)$ complete assignments
 - E.g., Boolean CSPs, including Boolean satisfiability (NP-complete)
- Discrete variables and infinite domains
 - E.g., job scheduling, variables are start/end times for each job
- Continuous variables
 - E.g., start/end times for Hubble Telescope observations



Varieties of constraints

- **Unary** constraints involve a single variable, e.g., $SA \neq green$
- **Binary** constraints involve pairs of variables, e.g., $SA \neq WA$
- **Higher-order** constraints involve 3 or more variables, e.g., cryptarithmic column constraints
- **Preferences** (soft constraints), e.g., *red* is better than *green* often representable by a cost for each variable assignment → **constrained optimization problems**

Real-world CSPs



- Assignment problems; e.g., who teaches what class?
- Timetabling problems; e.g., which class is offered when and where?
- Hardware configuration
- Spreadsheets
- Transportation scheduling
- Factory scheduling
- Floorplanning



Backtracking Search for CSPs

CSP as Standard Search Problem



Let's start with the straightforward, dumb approach, then fix it

- States are defined by the values assigned so far
 - **Initial state:** the empty assignment, \emptyset
 - **Successor function:** assign a value to an unassigned variable that does not conflict with current assignment \implies fail if no legal assignments (not fixable!)
 - **Goal test:** the current assignment is complete
1. This is the same for all CSPs!
 2. Every solution appears at depth n with n variables \implies use depth-first search
 3. Path is irrelevant, so can also use complete-state formulation
 4. Branching factor $b = (n - \ell)d$ at depth ℓ , hence $n!d^n$ leaves!

Backtracking search



- Variable assignments are **commutative**, i.e., $[WA = red \text{ then } NT = green]$ same as $[NT = green \text{ then } WA = red]$
- Only need to consider assignments to a single variable at each node $\implies b = d$ and there are d^n leaves
- Depth-first search for CSPs with single-variable assignments is called **backtracking** search
- Backtracking search is the basic uninformed algorithm for CSPs
- Can solve n -queens for $n \approx 25$

Algorithm



```

function BACKTRACKING-SEARCH(csp) returns a solution, or failure
    return BACKTRACK( $\emptyset$ , csp)
function BACKTRACK(assignment, csp) returns a solution, or failure
    if assignment is complete then return assignment
    var  $\leftarrow$  SELECT-UNASSIGNED-VARIABLE(csp)
    for each value in ORDER-DOMAIN-VALUES(var, assignment, csp) do
        if value is consistent with assignment then
            add {var = value} to assignment
            inferences  $\leftarrow$  INFERENCE(csp, var, value)
            if inferences  $\neq$  failure then
                add inferences to assignment
                result  $\leftarrow$  BACKTRACK(assignment, csp)
                if result  $\neq$  failure then
                    return result
            remove {var = value} and inferences from assignment
    return failure
    
```

Illustration

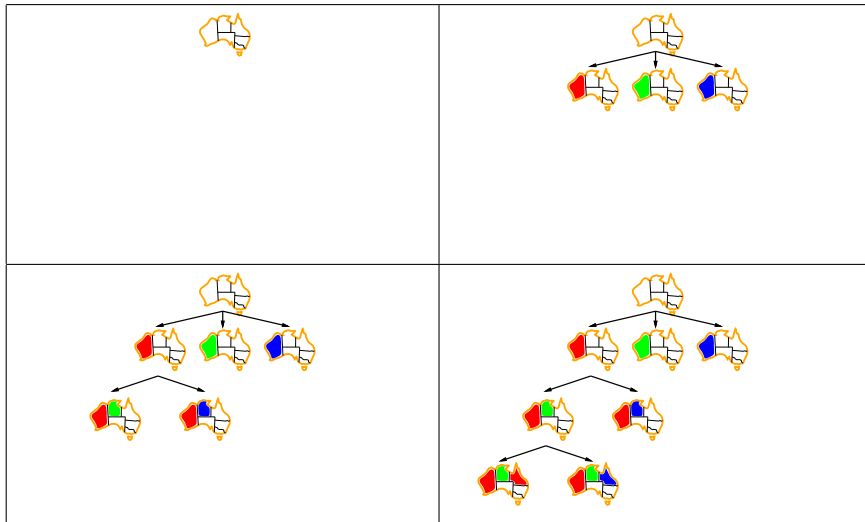


Backtracking Search for CSPs

Constraint
Propagation:
Inference in
CSPs

Local Search
for CSPs

The Structure
of Problems



Improving backtracking efficiency



General-purpose methods can give huge gains in speed:

1. Which variable should be assigned next?
2. In what order should its values be tried?
3. Can we detect inevitable failure early?
4. Can we take advantage of problem structure?



Minimum remaining values

The backtracking algorithm contains the line

$$var \leftarrow \text{SELECT-UNASSIGNED-VARIABLE}(csp)$$

We need a strategy for SELECT-UNASSIGNED-VARIABLE to choose the next unassigned variable in $\{X_1, X_2, \dots\}$


- **Minimum remaining values (MRV):** choose the variable with the fewest legal values
- MRV usually performs better than a random/static ordering, sometimes by a factor of 1,000 or more



Degree heuristic



- Tie-breaker among MRV variables
- **Degree heuristic:** choose the variable with the most constraints on remaining variables






Least constraining value

The backtracking algorithm contains the line

value in ORDER-DOMAIN-VALUES(*var*, *assignment*, *csp*)

- Given a variable, choose the least constraining value: the one that rules out the fewest values in the remaining variables



- Combining these heuristics makes 1000 queens feasible



Inference: Forward checking

- **Idea:** Keep track of remaining legal values for unassigned variables
→ Terminate search when any variable has no legal values


$$inferences \leftarrow \text{INFERENCE}(csp, var, value)$$

- **Note:** Forward checking propagates information from assigned to unassigned variables, but doesn't provide early detection for all failures!


Inference: Forward checking




a)




b)



c)



d)



Constraint Propagation: Inference in CSPs



Constraint propagation



- Reduce the number of legal values for a variable by using constraints → legal values for another variable also reduced
- Intertwined with search, or done as a preprocessing step
 - Sometimes the preprocessing can solve the whole problem!
- Enforcing local consistency in each part of a graph causes inconsistent values to be eliminated throughout the graph



Node consistency

- A single variable is **node-consistent** if all the values in the variable's domain satisfy the variable's **unary constraints**
 - South Australians dislike green,
 $SA = \{red, green, blue\} \rightarrow SA = \{red, blue\}$
- Eliminate all the unary constraints in a CSP by running node consistency



Arc consistency

- A variable in a CSP is arc-consistent if every value in its domain satisfies the variable's **binary constraints**. More formally,
$$X \rightarrow Y \text{ is consistent iff for every value } x \text{ of } X \text{ there is some allowed } y$$
- Run as a preprocessor before the search starts or after each assignment, must be run repeatedly until no inconsistency remains.
- Trade-off
 - Requires some overhead to do, but generally more effective than direct search
 - Eliminate large (inconsistent) parts of the state-space more effectively than search
- Need a systematic method for arc-checking
 - If Y loses a value, neighbors of Y need to be rechecked \rightarrow incoming arcs can become inconsistent again while outgoing arcs stay still

Algorithm

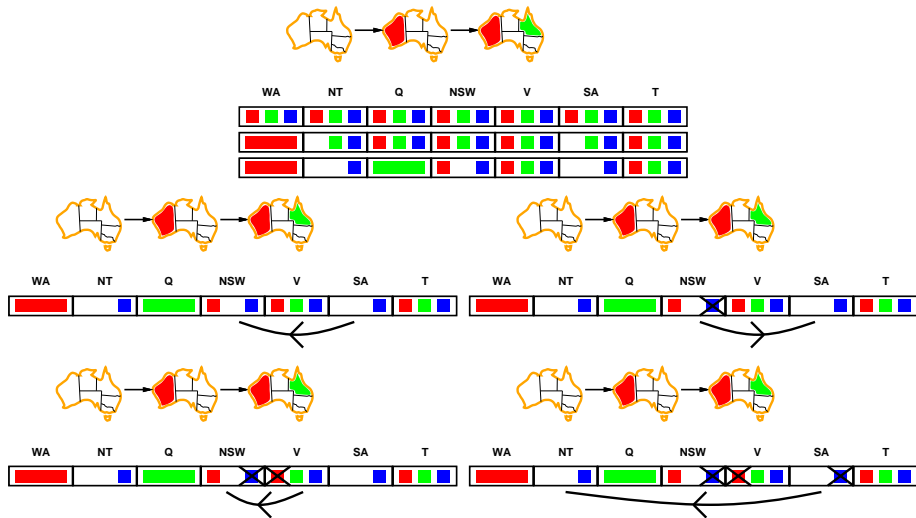


```

function AC-3(csp)
  returns false if an inconsistency is found and true otherwise
  inputs: csp, a binary CSP with components  $(\mathcal{X}, \mathcal{D}, \mathcal{C})$ 
  local variables: queue, a queue of arcs, initially all the arcs in csp
    while queue  $\neq \emptyset$  do
       $(X_i, X_j) \leftarrow \text{REMOVE-FIRST}(\textit{queue})$ 
      if REVISE(csp,  $X_i$ ,  $X_j$ ) then
        if size of  $D_i = 0$  then return false
        for each  $X_k$  in  $X_i.\text{NEIGHBORS} - \{X_j\}$  do
          add  $(X_k, X_i)$  to queue
    return true

function REVISE(csp,  $X_i$ ,  $X_j$ ) returns true iff we revise the domain of  $X_i$ 
  revised  $\leftarrow$  false
  for each  $x$  in  $D_i$  do
    if no value  $y$  in  $D_j$  allows  $(x, y)$  to satisfy
      the constraint between  $X_i$  and  $X_j$  then
      delete  $x$  from  $D_i$ 
      revised  $\leftarrow$  true
  return revised
  
```

Illustration





Local Search for CSPs

Local search for CSPs



- Complete-state formulation
 - The initial state assigns a value to every variable → violation
 - The search changes the value of one variable at a time → eliminate the violated constraints
- Min-conflicts heuristic: the minimum number of conflicts with other variables
- Min-conflicts is surprisingly effective for many CSPs.
 - Million-queens problem can be solved ~ 50 steps
 - Hubble Space Telescope: the time taken to schedule a week of observations down from 3 weeks to ~ 10 minutes

Algorithm



```

function MIN-CONFLICTS(csp, max_steps) returns a solution or failure
inputs: csp, a constraint satisfaction problem
        max_steps, the number of steps allowed before giving up
current  $\leftarrow$  an initial complete assignment for csp
for i = 1 to max_steps do
    if current is a solution for csp then return current
    var  $\leftarrow$  a randomly chosen conflicted variable from csp.VARIABLES
    value  $\leftarrow$  the value v for var that minimizes CONFLICTS(var, v, current, csp)
    set {var = value} in current
return failure
    
```



Performance of min-conflicts

- Given random initial state, can solve n -queens in almost constant time for arbitrary n with high probability (e.g., $n = 10,000,000$)
- The same appears to be true for any randomly-generated CSP *except* in a narrow range of the ratio

$$R = \frac{\text{number of constraints}}{\text{number of variables}}$$






Illustration: 8-queens




Figure 1: A two-step solution using min-conflicts for an 8-queens problem. At each stage, a queen is chosen for reassignment in its column. The number of conflicts (in this case, the number of attacking queens) is shown in each square. The algorithm moves the queen to the min-conflicts square, breaking ties randomly.

Improving min-conflicts



- The landscape of a CSP under the min-conflicts heuristic usually has a series of plateaux.
 - Millions of variable assignments that are only one conflict away from a solution
- **Plateau search:** allow sideways moves to another state with the same score
- **Tabu search:** keep a small list of recently visited states and forbid the algorithm to return to those states
- **Simulated annealing** can also be used

Constraint weighting



- Concentrate the search on the important constraints
- Each constraint is given a numeric weight, W_i , initially all 1.
- At each step, chooses a variable/value pair to change that has the lowest total weight of all violated constraints.
- Increase the weight of each constraint that is violated by the current assignment.

Local search in online setting



- Scheduling problems: online setting
 - A week's airline schedule may involve thousands of flights and tens of thousands of personnel assignments
 - The bad weather at one airport can render the schedule infeasible.
- The schedule should be repaired with a minimum number of changes.
 - Done easily with a local search starting from the current schedule
 - A backtracking search with the new set of constraints usually requires much more time and might find a solution with many changes from the current schedule



The Structure of Problems

Independent subproblems



- A problem CSP is decomposed into **independent subproblems** CSP_i ;
 - Independent subproblems are identifiable as connected components of constraint graph




Figure 2: Tasmania and mainland are **independent subproblems**

Independent subproblems (cont.)



- If assignment S_i is a solution of CSP_i then

$$S = \bigcup_{i=1} S_i \text{ is a solution of } CSP$$

- Suppose a problem of n variables can be broken into independent subproblems of only c variables
 - Worst-case solution cost is $O(\frac{n}{c}d^c)$
 - Original CSP has worst-case solution $O(d^n)$

Tree-structured CSPs



- A CSP is a tree-structured CSP if the constraint graph has no loops

Theorem 1

A tree-structured CSP can be solved in $O(n \times d^2)$ time, compare to general CSPs, where worst-case time is $O(d^n)$

Proof

self-exercise ■



Tree-structured CSPs (cont.)

- Using `TOPOLOGICALSORT` to create an ordering of the variables such that each variable appears after its parent in the tree




Figure 3: (a) The constraint graph of a tree-structured CSP. (b) A linear ordering of the variables consistent with the tree with *A* as the root. This is known as a **topological sort** of the variables.

Algorithm



```

function TREE-CSP-SOLVER(csp) returns a solution, or failure
inputs: csp, a CSP with components  $\mathcal{X}$ ,  $\mathcal{D}$ ,  $\mathcal{C}$ 
     $n \leftarrow |\mathcal{X}|$ 
    assignment  $\leftarrow \emptyset$ 
    root  $\leftarrow$  any variable in  $\mathcal{X}$ 
     $\mathcal{X} \leftarrow \text{TOPOLOGICALSORT}(\mathcal{X}, \textit{root})$ 
    for  $j = n$  down to 2 do
        MAKE-ARC-CONSISTENT(PARENT( $X_j$ ),  $X_j$ )
        if it cannot be made consistent then return failure
    for  $i = 1$  to  $n$  do
        assignment[ $X_i$ ]  $\leftarrow$  any consistent value from  $D_i$ 
        if there is no consistent value then return failure
    return assignment
    
```

Nearly tree-structured CSPs



- **Cycle cutset:** a set of variables such that the remaining constraint graph is a tree




Figure 4: The removal of *SA* makes the constraint graph to be a tree



Cutset Conditioning


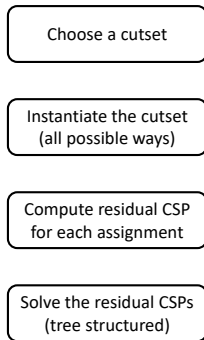
Defining
Constraint
Satisfaction
Problems

Backtracking
Search for
CSPs

Constraint
Propagation:
Inference in
CSPs

Local Search
for CSPs

The Structure
of Problems



- If the cycle cutset has size c , then the total run time is $O(d^c(n - c)d^2)$

Tree Decomposition



- **Idea:** create a tree-structured graph of mega-variables
 - Each mega-variable encodes part of the original CSP
 - Subproblems overlap to ensure consistent solutions
- A **tree decomposition** must satisfy the following three requirements:
 1. Every variable in the original problem appears in at least one of the subproblems.
 2. If two variables are connected by a constraint in the original problem, they must appear together (along with the constraint) in at least one of the subproblems.
 3. If a variable appears in two subproblems in the tree, it must appear in every subproblem along the path connecting those subproblems.

Tree Decomposition (cont.)




Figure 5: A tree decomposition of the constraint graph

The structure of values



- Consider the map-coloring problem with n colors.
- For every consistent solution, there is actually a set of $n!$ solutions formed by permuting the color names.
 - E.g., WA , NT , and SA must all have different colors, but there are $3!$ ways to assign the three colors to these three regions.
- **Symmetry-breaking constraint:** Impose an arbitrary ordering constraint that requires the values to be in alphabetical order.
 - E.g., $NT < SA < WA \rightarrow$ only one solution possible $\{NT = \text{blue}, SA = \text{green}, WA = \text{red}\}$

References



Goodfellow, I., Bengio, Y., and Courville, A. (2016).

Deep learning.

MIT press.



Nguyen, T. (2018).

Artificial intelligence slides.

Technical report, HCMC University of Sciences.



Russell, S. and Norvig, P. (2016).

Artificial intelligence: a modern approach.

Pearson Education Limited.